

```

class Discriminator(nn.Module):
    ##TODO1 implement the discriminator (critic)
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(in_channels=1, out_channels=128, kernel_size=5, stride=2, padding=2)
        self.conv2 = nn.Conv2d(in_channels=128, out_channels=256, kernel_size=5, stride=2, padding=2)
        self.conv3 = nn.Conv2d(in_channels=256, out_channels=512, kernel_size=5, stride=2, padding=2)

        self.fc = nn.Linear(in_features=512*4*4, out_features=1)

    def forward(self, x):
        x = F.leaky_relu(self.conv1(x), 0.1)
        x = F.leaky_relu(self.conv2(x), 0.1)
        x = F.leaky_relu(self.conv3(x), 0.1)

        x = x.view(-1, 512*4*4)
        x = self.fc(x)

        return x

```

```

class Generator(nn.Module):
    ##TODO2 implement the generator (actor)
    def __init__(self):
        super().__init__()

        self.fc = nn.Linear(128, 512*4*4)

        self.conv_block1 = nn.Sequential(
            nn.ConvTranspose2d(512, 256, kernel_size=5, stride=2, padding=2, output_padding=1),
            nn.BatchNorm2d(256),
            nn.ReLU(inplace=True),
        )
        self.conv_block2 = nn.Sequential(
            nn.ConvTranspose2d(256, 128, kernel_size=5, stride=2, padding=2, output_padding=1),
            nn.BatchNorm2d(128),
            nn.ReLU(inplace=True),
        )
        self.conv_block3 = nn.Sequential(
            nn.ConvTranspose2d(128, 1, kernel_size=5, stride=2, padding=2, output_padding=1),
            nn.Sigmoid(),
        )

    def forward(self, x):
        x = self.fc(x)
        x = x.view(-1, 512, 4, 4)

        x = self.conv_block1(x)
        x = self.conv_block2(x)
        x = self.conv_block3(x)

        x = F.interpolate(x, size=(28, 28), mode='bilinear', align_corners=False)

        return x

```

TODO 3: What is the input and output shape of the generator and discriminator network? Verify that the implemented networks are the same as the answer you have provided.

```

from torchinfo import summary
# Test the discriminator
print(summary(discriminator, input_size=((1, 1, 28, 28))))

# Test the generator
print(summary(generator, input_size=((1, 128))))

```

```

=====
Layer (type:depth-idx)          Output Shape       Param #
=====
Discriminator                  [1, 1]           --
|---Conv2d: 1-1                [1, 128, 14, 14]   3,328
|---Conv2d: 1-2                [1, 256, 7, 7]    819,456
|---Conv2d: 1-3                [1, 512, 4, 4]   3,277,312
|---Linear: 1-4                [1, 1]           8,193
=====

Total params: 4,108,289
Trainable params: 4,108,289
Non-trainable params: 0
Total mult-adds (M): 93.25
=====

Input size (MB): 0.00
Forward/backward pass size (MB): 0.37
Params size (MB): 16.43
Estimated Total Size (MB): 16.80
=====
```

Layer (type:depth-idx)	Output Shape	Param #
Generator	[1, 1, 28, 28]	1,056,768
---Linear: 1-1	[1, 8192]	--
---Sequential: 1-2	[1, 256, 8, 8]	--
---ConvTranspose2d: 2-1	[1, 256, 8, 8]	3,277,056
---BatchNorm2d: 2-2	[1, 256, 8, 8]	512
---ReLU: 2-3	[1, 256, 8, 8]	--
---Sequential: 1-3	[1, 128, 16, 16]	--
---ConvTranspose2d: 2-4	[1, 128, 16, 16]	819,328
---BatchNorm2d: 2-5	[1, 128, 16, 16]	256
---ReLU: 2-6	[1, 128, 16, 16]	--
---Sequential: 1-4	[1, 1, 32, 32]	--
---ConvTranspose2d: 2-7	[1, 1, 32, 32]	3,201
---Sigmoid: 2-8	[1, 1, 32, 32]	--

```

=====
Total params: 5,157,121
Trainable params: 5,157,121
Non-trainable params: 0
Total mult-adds (M): 423.81
=====

Input size (MB): 0.00
Forward/backward pass size (MB): 0.86
Params size (MB): 20.63
Estimated Total Size (MB): 21.49
=====
```

```
NUM_ITERATION = 3000
BATCH_SIZE = 32
fixed_z = torch.randn(8, 128).cuda()
def schedule(i):
    lr = 1e-4
    if(i > 2500): lr *= 0.1
    return lr
losses = {'D' : [None], 'G' : [None]}

## TODO4 initialize missing hyperparameter and optimzer
G_optimizer = torch.optim.Adam(generator.parameters(), lr=1e-4, betas=(0.0, 0.9))
D_optimizer = torch.optim.Adam(discriminator.parameters(), lr=1e-4, betas=(0.0, 0.9))
GP_lambda = 10
n_critic = 5
```

```
# TODO5 implement dataloader
import torch
from torch.utils.data import DataLoader, Dataset

class CustomDataset(Dataset):
    def __init__(self, data, z_dim):
        self.data = data
        self.z_dim = z_dim

    def __len__(self):
        return len(self.data)

    def __getitem__(self, idx):
        x = torch.FloatTensor(self.data[idx])
        z = torch.randn(self.z_dim)
        epsilon = torch.rand(1)

        return x, z, epsilon

Z_DIM = 128
BATCH_SIZE = 32

custom_dataset = CustomDataset(trainX, Z_DIM)
dataloader = DataLoader(custom_dataset, batch_size=BATCH_SIZE, shuffle=True)
```

```

from tqdm import tqdm
for i in tqdm(range(NUM_ITERATION)):
    ## TODO6 update learning rate
    for params in D_optimizer.param_groups:
        params['lr'] = schedule(i)
    for params in G_optimizer.param_groups:
        params['lr'] = schedule(i)

    discriminator.train()
    for t in range(n_critic):
        ## TODO7 line 4: sample data from dataloader
        real_images, z, epsilon = next(iter(dataloader))

        real_images = real_images.cuda()
        z = z.cuda()
        epsilon = epsilon.cuda().view(-1, 1, 1, 1)

        ## TODO8 line5-7 : calculate discriminator loss
        D_optimizer.zero_grad()

        fake_images = generator(z)

        alpha = epsilon.view(-1, 1, 1, 1)
        interpolated_images = alpha * real_images + (1 - alpha) * fake_images.detach()
        interpolated_images.requires_grad_(True)

        real_predict = discriminator(real_images)
        fake_predict = discriminator(fake_images.detach())
        interpolated_predict = discriminator(interpolated_images)

        gradients = torch.autograd.grad(outputs=interpolated_predict, inputs=interpolated_images,
                                         grad_outputs=torch.ones_like(interpolated_predict),
                                         create_graph=True, retain_graph=True)[0]
        gradients_penalty = ((gradients.norm(2, dim=1) - 1) ** 2).mean() * GP_lambda

        d_loss = fake_predict.mean() - real_predict.mean() + gradients_penalty
        losses['D'].append(d_loss.item())

        ## TODO9 : line 9 update discriminator loss
        d_loss.backward()
        D_optimizer.step()

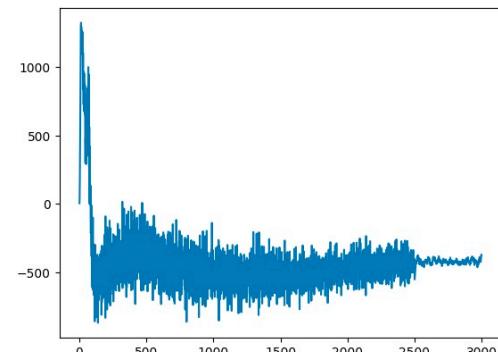
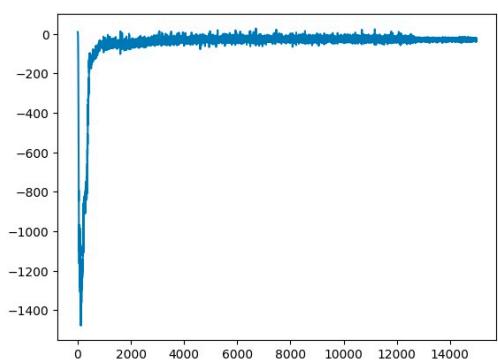
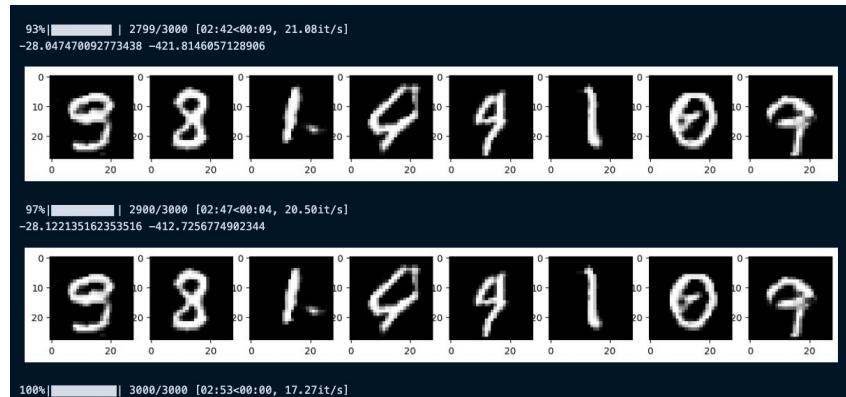
    ## TODO10 : line 11-12 calculate and update the generator loss
    generator.train()
    G_optimizer.zero_grad()

    fake_images = generator(z)
    fake_predict = discriminator(fake_images)
    g_loss = -fake_predict.mean()
    losses['G'].append(g_loss.item())

    g_loss.backward()
    G_optimizer.step()

# Output visualization : If your reimplementation is correct, the generated images should start resembling a digit character after 500 iterations.
discriminator.eval()
generator.eval()
if(i % 100 == 0):
    plt.figure(figsize = (15,75))
    print(losses['D'][-1], losses['G'][-1])
    with torch.no_grad():
        res = generator(fixed_z).cpu().detach().numpy()
    for k in range(8):
        plt.subplot(int('18{}'.format(k+1)))
        plt.imshow(res[k].transpose(1, 2, 0)[..., 0], cmap = 'gray' )
    plt.show()

```



```

class ConvBlock(nn.Module):
    def __init__(self, in_channels, out_channels, kernel_size=4, stride=2, padding=1, with_norm=True):
        super().__init__()
        self.conv = nn.Conv2d(in_channels=in_channels, out_channels=out_channels, kernel_size=kernel_size, stride=stride, padding=padding)
        self.instancenorm = nn.InstanceNorm2d(out_channels)
        self.activation = nn.LeakyReLU(0.2)
        self.with_norm = with_norm

    def forward(self, x):
        x = self.conv(x)
        if self.with_norm:
            x = self.instancenorm(x)
        x = self.activation(x)

        return x

class Discriminator(nn.Module):
    #TODO011 implement the discriminator network
    def __init__(self):
        super().__init__()
        self.block64 = ConvBlock(in_channels=6, out_channels=64, with_norm=False)
        self.block128 = ConvBlock(in_channels=64, out_channels=128)
        self.block256 = ConvBlock(in_channels=128, out_channels=256)
        self.block512 = ConvBlock(in_channels=256, out_channels=512, stride=1)

        self.conv = nn.Conv2d(in_channels=512, out_channels=1, kernel_size=4, stride=1, padding=1)
        self.sigmoid = nn.Sigmoid()

    def forward(self, x):
        x = self.block64(x)
        x = self.block128(x)
        x = self.block256(x)
        x = self.block512(x)
        x = self.conv(x)
        x = self.sigmoid(x)

        return x

discriminator = Discriminator().cuda()

```

```

#TODO012 verify the discriminator
from torchinfo import summary
summary(discriminator, input_size=(1, 6, 256, 256))

```

Layer (type:depth-idx)	Output Shape	Param #
Discriminator	[1, 1, 30, 30]	--
—Conv2d: 1-1	[1, 64, 128, 128]	--
—LeakyReLU: 2-1	[1, 64, 128, 128]	6,208
—ConvBlock: 1-2	[1, 128, 64, 64]	--
—Conv2d: 2-3	[1, 128, 64, 64]	131,200
—InstanceNorm2d: 2-4	[1, 128, 64, 64]	--
—LeakyReLU: 2-5	[1, 128, 64, 64]	--
—ConvBlock: 1-3	[1, 256, 32, 32]	--
—Conv2d: 2-6	[1, 256, 32, 32]	524,544
—InstanceNorm2d: 2-7	[1, 256, 32, 32]	--
—LeakyReLU: 2-8	[1, 256, 32, 32]	--
—ConvBlock: 1-4	[1, 512, 31, 31]	--
—Conv2d: 2-9	[1, 512, 31, 31]	2,097,664
—InstanceNorm2d: 2-10	[1, 512, 31, 31]	--
—LeakyReLU: 2-11	[1, 512, 31, 31]	--
—Conv2d: 1-5	[1, 1, 30, 30]	8,193
—Sigmoid: 1-6	[1, 1, 30, 30]	--

Total params: 2,767,809  
Trainable params: 2,767,809  
Non-trainable params: 0  
Total mult-adds (G): 3.20

Input size (MB): 1.57  
Forward/backward pass size (MB): 18.62  
Params size (MB): 11.07  
Estimated Total Size (MB): 31.27

```

class ConvTBlock(nn.Module):
    def __init__(self, in_channels, out_channels, kernel_size=4, stride=2, padding=1, p=0.5, with_norm=True, with_relu=True):
        super(ConvTBlock, self).__init__()
        self.conv = nn.ConvTranspose2d(in_channels=in_channels, out_channels=out_channels, kernel_size=kernel_size, stride=stride, padding=padding)
        self.instancenorm = nn.InstanceNorm2d(out_channels)
        self.dropout = nn.Dropout(p=p)
        self.relu = nn.ReLU()
        self.with_relu = with_relu
        self.with_norm = with_norm

    def forward(self, x, x_enc = None):
        if x_enc is not None:
            x = torch.cat((x, x_enc), 1)
        x = self.conv(x)
        if self.with_norm:
            x = self.instancenorm(x)
        x = self.dropout(x)
        if self.with_relu:
            x = self.relu(x)
        return x

```

```

class Generator(nn.Module):
    #TODO13 implement the generator network
    def __init__(self):
        super().__init__()

        # encoder
        self.encoder1 = ConvBlock(in_channels=3, out_channels=64, with_norm=False)
        self.encoder2 = ConvBlock(in_channels=64, out_channels=128)
        self.encoder3 = ConvBlock(in_channels=128, out_channels=256)
        self.encoder4 = ConvBlock(in_channels=256, out_channels=512)
        self.encoder5 = ConvBlock(in_channels=512, out_channels=512)
        self.encoder6 = ConvBlock(in_channels=512, out_channels=512)
        self.encoder7 = ConvBlock(in_channels=512, out_channels=512)
        self.encoder8 = ConvBlock(in_channels=512, out_channels=512, with_norm=False)

        # decoder
        self.decoder1 = ConvTBlock(in_channels=512, out_channels=512)
        self.decoder2 = ConvTBlock(in_channels=1024, out_channels=512)
        self.decoder3 = ConvTBlock(in_channels=1024, out_channels=512)
        self.decoder4 = ConvTBlock(in_channels=1024, out_channels=512)
        self.decoder5 = ConvTBlock(in_channels=1024, out_channels=256)
        self.decoder6 = ConvTBlock(in_channels=512, out_channels=128)
        self.decoder7 = ConvTBlock(in_channels=256, out_channels=64)
        self.decoder8 = ConvTBlock(in_channels=128, out_channels=3, with_norm=False, with_relu=False)

        self.tanh = nn.Tanh()

    def forward(self, x):
        e1 = self.encoder1(x)
        e2 = self.encoder2(e1)
        e3 = self.encoder3(e2)
        e4 = self.encoder4(e3)
        e5 = self.encoder5(e4)
        e6 = self.encoder6(e5)
        e7 = self.encoder7(e6)
        e8 = self.encoder8(e7)

        d1 = self.decoder1(e8)
        d2 = self.decoder2(d1, e7)
        d3 = self.decoder3(d2, e6)
        d4 = self.decoder4(d3, e5)
        d5 = self.decoder5(d4, e4)
        d6 = self.decoder6(d5, e3)
        d7 = self.decoder7(d6, e2)
        d8 = self.decoder8(d7, e1)

        o = self.tanh(d8)
        return o

```

```
generator = Generator().cuda()
```

```

#TODO14 verify the generator
from torchinfo import summary
summary(generator, input_size=(1, 3, 256, 256))

```

Layer (type:depth-idx)	Output Shape	Param #
Generator	[1, 3, 256, 256]	--
ConvBlock: 1-1	[1, 64, 128, 128]	--
└Conv2d: 2-1	[1, 64, 128, 128]	3,136
└LeakyReLU: 2-2	[1, 64, 128, 128]	--
ConvBlock: 1-2	[1, 128, 64, 64]	--
└Conv2d: 2-3	[1, 128, 64, 64]	131,200
└InstanceNorm2d: 2-4	[1, 128, 64, 64]	--
└LeakyReLU: 2-5	[1, 128, 64, 64]	--
ConvBlock: 1-3	[1, 256, 32, 32]	--
└Conv2d: 2-6	[1, 256, 32, 32]	524,544
└InstanceNorm2d: 2-7	[1, 256, 32, 32]	--
└LeakyReLU: 2-8	[1, 256, 32, 32]	--
ConvBlock: 1-4	[1, 512, 16, 16]	--
└Conv2d: 2-9	[1, 512, 16, 16]	2,097,664
└InstanceNorm2d: 2-10	[1, 512, 16, 16]	--
└LeakyReLU: 2-11	[1, 512, 16, 16]	--
ConvBlock: 1-5	[1, 512, 8, 8]	--
└Conv2d: 2-12	[1, 512, 8, 8]	4,194,816
└InstanceNorm2d: 2-13	[1, 512, 8, 8]	--
└LeakyReLU: 2-14	[1, 512, 8, 8]	--
ConvBlock: 1-6	[1, 512, 4, 4]	--
└Conv2d: 2-15	[1, 512, 4, 4]	4,194,816
└InstanceNorm2d: 2-16	[1, 512, 4, 4]	--
└LeakyReLU: 2-17	[1, 512, 4, 4]	--
ConvBlock: 1-7	[1, 512, 2, 2]	--
└Conv2d: 2-18	[1, 512, 2, 2]	4,194,816
└InstanceNorm2d: 2-19	[1, 512, 2, 2]	--
└LeakyReLU: 2-20	[1, 512, 2, 2]	--
ConvBlock: 1-8	[1, 512, 1, 1]	--
└Conv2d: 2-21	[1, 512, 1, 1]	4,194,816
└LeakyReLU: 2-22	[1, 512, 1, 1]	--
ConvTBlock: 1-9	[1, 512, 2, 2]	--
└ConvTranspose2d: 2-23	[1, 512, 2, 2]	4,194,816
└InstanceNorm2d: 2-24	[1, 512, 2, 2]	--
└Dropout: 2-25	[1, 512, 2, 2]	--
└ReLU: 2-26	[1, 512, 2, 2]	--
ConvTBlock: 1-10	[1, 512, 4, 4]	--
└ConvTranspose2d: 2-27	[1, 512, 4, 4]	8,389,120
└InstanceNorm2d: 2-28	[1, 512, 4, 4]	--
└Dropout: 2-29	[1, 512, 4, 4]	--
└ReLU: 2-30	[1, 512, 4, 4]	--
ConvTBlock: 1-11	[1, 512, 8, 8]	--
└ConvTranspose2d: 2-31	[1, 512, 8, 8]	8,389,120
└InstanceNorm2d: 2-32	[1, 512, 8, 8]	--
└Dropout: 2-33	[1, 512, 8, 8]	--
└ReLU: 2-34	[1, 512, 8, 8]	--
ConvTBlock: 1-12	[1, 512, 16, 16]	--
└ConvTranspose2d: 2-35	[1, 512, 16, 16]	8,389,120
└InstanceNorm2d: 2-36	[1, 512, 16, 16]	--
└Dropout: 2-37	[1, 512, 16, 16]	--
└ReLU: 2-38	[1, 512, 16, 16]	--
ConvTBlock: 1-13	[1, 256, 32, 32]	--
└ConvTranspose2d: 2-39	[1, 256, 32, 32]	4,194,560
└InstanceNorm2d: 2-40	[1, 256, 32, 32]	--
└Dropout: 2-41	[1, 256, 32, 32]	--
└ReLU: 2-42	[1, 256, 32, 32]	--
ConvTBlock: 1-14	[1, 128, 64, 64]	--
└ConvTranspose2d: 2-43	[1, 128, 64, 64]	1,048,704
└InstanceNorm2d: 2-44	[1, 128, 64, 64]	--
└Dropout: 2-45	[1, 128, 64, 64]	--
└ReLU: 2-46	[1, 128, 64, 64]	--
ConvTBlock: 1-15	[1, 64, 128, 128]	--
└ConvTranspose2d: 2-47	[1, 64, 128, 128]	262,288
└InstanceNorm2d: 2-48	[1, 64, 128, 128]	--
└Dropout: 2-49	[1, 64, 128, 128]	--
└ReLU: 2-50	[1, 64, 128, 128]	--
ConvTBlock: 1-16	[1, 3, 256, 256]	--
└ConvTranspose2d: 2-51	[1, 3, 256, 256]	6,147
└Dropout: 2-52	[1, 3, 256, 256]	--
Tanh: 1-17	[1, 3, 256, 256]	--
<hr/>		
Total params: 54,409,603		
Trainable params: 54,409,603		
Non-trainable params: 0		
Total mult-adds (G): 18.14		
<hr/>		
Input size (MB): 0.79		
Forward/backward pass size (MB): 33.72		
Params size (MB): 217.64		
Estimated Total Size (MB): 252.15		
<hr/>		

```

# TOD015 implement a dataloader
from torchvision import transforms
from torch.utils.data import DataLoader, Dataset, RandomSampler

class Pix2PixDataset(Dataset):
    def __init__(self, x, y, transformer):
        self.x = x.astype(np.float32)
        self.y = y.astype(np.float32)
        self.transformer = transformer

    def __getitem__(self, index):
        x = torch.tensor(self.x[index])
        y = torch.tensor(self.y[index])

        if self.transformer:
            xy_combined = torch.cat((x.unsqueeze(0), y.unsqueeze(0)), dim=0)
            xy_transformed = self.transformer(xy_combined)

            x_transformed, y_transformed = torch.split(xy_transformed, [1, 1], dim=0)
            x_transformed = x_transformed.squeeze(0)
            y_transformed = y_transformed.squeeze(0)
        else:
            x_transformed, y_transformed = x, y

        return x_transformed, y_transformed

    def __len__(self):
        return self.x.shape[0]

transformer = transforms.Compose([
    transforms.Resize(286),
    transforms.RandomCrop(256),
    transforms.RandomHorizontalFlip(0.5),
])

train_dataset = Pix2PixDataset(trainX, trainY, transformer)
train_loader = DataLoader(train_dataset, batch_size=1, shuffle=True)

```

```

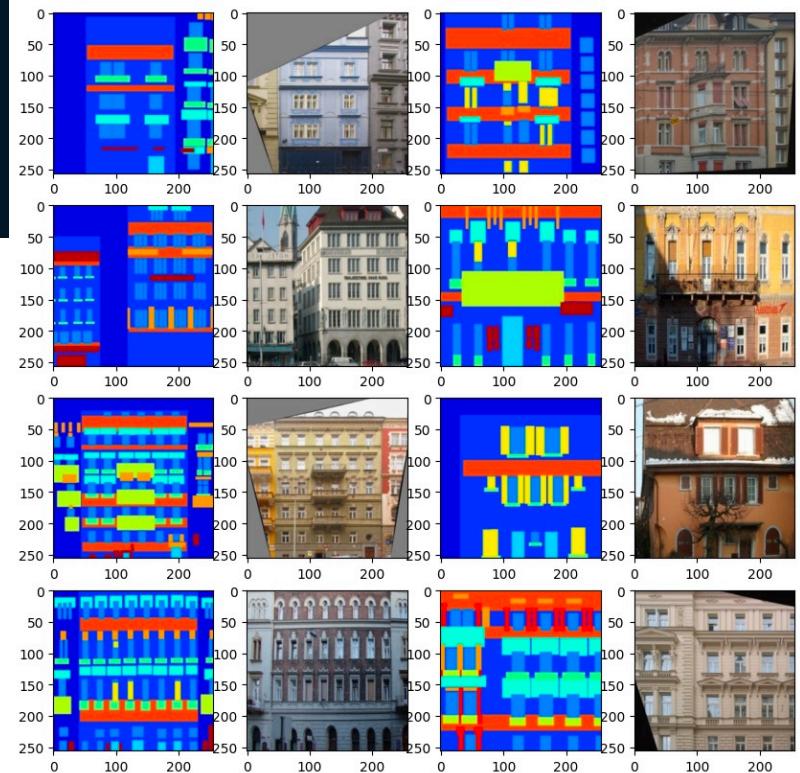
# TOD016 show that the dataloader is working properly
def to_image(x):
    x = x.numpy()
    return (0.5 * x.transpose((1, 2, 0)) + 0.5)[..., ::-1]

```

```

plt.figure(figsize = (10, 10))
for i in range(8):
    x, y = next(iter(train_loader))
    plt.subplot(4, 4, 2*i+1)
    plt.imshow(to_image(x[0]), cmap="gray")
    plt.subplot(4, 4, 2*i+2)
    plt.imshow(to_image(y[0]), cmap="gray")
    plt.show()

```



```

losses = {'D' : [None], 'G' : [None]}

for i in tqdm(range(20001)):
    #TODO017 sample the data from the dataloader
    x, y = next(iter(train_loader))
    x = x.cuda()
    y = y.cuda()

    #TODO018 calculate the discriminator loss and update the discriminator
    discriminator.train()
    D_optimizer.zero_grad()

    # Real pairs
    real_pairs = torch.cat((x, y), dim=1)
    D_real = discriminator(real_pairs)
    D_real_loss = torch.log(D_real + 1e-12).mean()

    # Fake pairs
    fake_y = generator(x)
    fake_pairs = torch.cat((x, fake_y), dim=1)
    D_fake = discriminator(fake_pairs.detach())
    D_fake_loss = torch.log(1 - D_fake + 1e-12).mean()

    # Combined D loss
    D_loss = -0.5 * (D_real_loss + D_fake_loss) # Multiply by -0.5 as suggested
    D_loss.backward()
    D_optimizer.step()

    losses["D"].append(D_loss.item())

    #TODO019 calculate the generator loss and update the generator
    generator.train()
    G_optimizer.zero_grad()

    G_fake = discriminator(fake_pairs)
    G_gan_loss = -torch.log(G_fake + 1e-12).mean()

    G_L1_loss = F.l1_loss(fake_y, y) * LAMBDA

    G_loss = G_gan_loss + G_L1_loss
    G_loss.backward()
    G_optimizer.step()

    losses["G"].append(G_loss.item())

# Output visualization : If your reimplementation is correct, the generated images should start resembling a facade after 2,500 iterations
discriminator.eval()
generator.eval()
if i % 2500 == 0:
    with torch.no_grad():
        print(losses['D'][-1], losses['G'][-1])
        plt.figure(figsize = (40, 16))
        gs1 = gridspec.GridSpec(4, 4)
        gs1.update(wspace=0.025)

        sampleX_vis = 0.5 * valX[:16][:, ::-1, :, :] + 0.5
        sampleY = 0.5 * valY[:16][:, ::-1, :, :] + 0.5
        sampleX = torch.tensor(valX[:16]).cuda()
        pred_val = 0.5 * generator(sampleX.cpu().detach().numpy())[:, ::-1, :, :] + 0.5
        vis = np.concatenate([sampleX_vis, sampleY, pred_val], axis = 3)
        for i in range(vis.shape[0]):
            ax1 = plt.subplot(gs1[i])
            plt.title('Input / GT / predicted')
            plt.axis('off')
            plt.imshow(vis[i].transpose(1, 2, 0), cmap="gray" )
        plt.show()

```

